



TAKE-HOME ASSIGNMENT

Product Manager

Core Platform · Copilot

4-day window · ~8 hours of work · AI use encouraged

About this assignment

Superleap is an AI-native CRM used by sales teams of 50 to 5,000 people across education, healthcare, BFSI, real estate, tech marketplaces, etc. The core platform team builds the capabilities that everyone else uses: our deployment team configures them for each client, and our sales team sells them.

This assignment is built around a real kind of problem at that level: an AI capability that is technically impressive but not yet trusted or used enough.

- **Time window:** 4 days from receipt. Plan for about 8 hours of work.
- **Use AI as much as you want.** We expect you to. Include a short log of the tools and key prompts you used.
- **The data below is illustrative** but realistic. Make and state assumptions where you need to.
- **We value depth over coverage.**

The situation: Superleap Copilot

We launched Superleap Copilot sometime back, a natural-language assistant inside the CRM. Users open it from a sidebar and type or speak a request. It can:

- answer questions about pipeline and performance ("How many hot leads did Pune get this week?"),
- find records ("Show me patients who missed their second consultation"),
- draft messages for WhatsApp, SMS and email,
- update records ("Mark this lead as not interested, reason: fees").

It runs on an in-house model orchestrator which acts as a gateway to third-party large language models, with our own layer that reads each customer's CRM setup. Every customer's setup is different: their own objects, fields, stages and terminology.

Internal demos are impressive and it helps us win deals. But usage after go-live is not where it should be. The founders' question to you is simple: **"Is Copilot a product people rely on, or a demo feature? Fix it."**

The data

This covers the last 90 days across 25 customers and about 12,000 users with Copilot access.

Usage

Metric	Value
Weekly active Copilot users	14% of users with access
Weekly active, managers	41%
Weekly active, frontline (reps, counsellors, agents)	9%
Requests per active user per week	6
Opened from the sidebar button	81% of sessions
Opened from an in-context suggestion	4% of sessions

What people ask for

Request type	Share of requests	Thumbs-down rate
Pipeline and performance questions	52%	27%
Find records	21%	14%
Draft messages	14%	8%
Update records	8%	11%
Other	5%	22%

Quality, speed and cost

Metric	Value
Overall thumbs-down rate	19%
Requests in Hinglish or a regional language	31% of requests, 34% thumbs-down
Requests that touch customer-specific fields or objects	44% of requests, 29% thumbs-down
Response time	4.8 seconds median, 14 seconds at the 95th percentile
Model cost per request	₹1.9 on average, ₹6.5 for pipeline and performance questions
Users who gave a thumbs-down and used Copilot again within 2 weeks	23%

What users say

- *"The number didn't match my dashboard. After that I stopped asking."*
- *"It doesn't know what a 'hot lead' means for us."*
- *"It's faster to just click the filters."*
- *"I don't trust it to update records without checking."*
- *"Great for my weekly review. I don't need it during calls."* (sales manager)
- *"Where is it? I forgot it existed."* (counsellor)

What we want from you

There are four parts. A is diagnosis, B is AI quality, C is the next-generation bet, and D is presenting it.

Part A: Diagnose (1–2 pages)

1. Answer the business's question. Is Copilot a product people rely on, or a demo feature? Why?
2. What is really holding adoption back? Separate the causes into kinds, for example capability, trust, discoverability, workflow fit, or cost. Rank them.
3. Is "weekly active Copilot users" the right metric? If not, what should we measure instead?
4. What would you stop doing, and what data would you want before deciding anything big?

Part B: Design the evals (1–2 pages, plus a small eval run)

Pick the one request type from the data where better quality would matter most.

1. Design an evaluation suite for it. Cover:
 - how you build the test set, given that every customer's CRM setup is different,
 - what you measure and how you grade it (automated checks, AI grading, human review),
 - the quality bar a change must clear before it ships, and how you catch regressions.
 2. What do public AI benchmarks for agents and CRM tasks measure, and what do they miss for a product like ours?
 3. **Run a small version.** Write 20–30 test cases, run them against any model or tool you like, grade the results and share what you found. We care more about what you learned than the score.
-

Part C: The next-generation bet (2 pages max, plus a prototype)

Propose one platform capability that would make Superleap's AI something users rely on every day, 18–24 months from now. It must be a platform capability, not a feature for one industry.

1. What is it, who is it for, and why now?
2. How does our deployment team configure it for a new customer, and how does our sales team sell it?
3. What are the hard trade-offs: accuracy, speed, cost per user, build versus buy? Show rough economics.
4. How will you know it's working in 90 days?
5. Build a prototype of the core experience. It can be clickable or working. Show how it behaves for a real user in a real workflow.

Part D: 5-minute Loom

Present your recommendation to our founders: what you would do with Copilot in the next 90 days, and your bet for the next 2 years. Keep your face on camera.

How to submit

Send one PDF, doc, or Notion link covering Parts A, B and C, plus links to your eval run (Part B), prototype and prompt log (Part C), and Loom (Part D).

- Keep the written parts under 6 pages.
 - Make sure every link you share is viewable without requesting access.
 - Submit within 4 days of receiving this assignment.
-